

# *Database Management Systems*

## *Chapter 1*



Instructor: Rick Schumeyer  
richard.schumeyer@villanova.edu

## *What Is a DBMS?*



- ❖ A very large, integrated collection of data.
- ❖ Models real-world *enterprise*.
  - Entities (e.g., students, courses)
  - Relationships (e.g., Madonna is taking CS564)
- ❖ A *Database Management System (DBMS)* is a software package designed to store and manage databases.

## *Brief History*

- ❖ early 1960s: first general purpose database by Charles Bachman from GE. Used the network data model.
- ❖ late 1960s: IBM developed Information Management System (IMS). Used the hierarchical data model. Led to SABRE, the airline reservation system developed by AA and IBM. Still in use today.

## *Brief History*

- ❖ 1970: Edgar Codd of IBM developed the relational data model. Led to several DBMS based on relational model, as well as important theoretical results. Codd wins Turing award.
- ❖ 1980s: relational model dominant. SQL standard.
- ❖ Late 1980s, 1990s: DBMS vendors extend systems, allowing more complex data types (images, text).

## *Files vs. DBMS*

- ❖ Application must stage large datasets between main memory and secondary storage (e.g., buffering, page-oriented access, 32-bit addressing, etc.)
- ❖ Special code for different queries
- ❖ Must protect data from inconsistency due to multiple concurrent users
- ❖ Crash recovery
- ❖ Security and access control

## *Why Use a DBMS?*



- ❖ Data independence and efficient access.
- ❖ Reduced application development time.
- ❖ Data integrity and security.
- ❖ Uniform data administration.
- ❖ Concurrent access, recovery from crashes.

## Why Study Databases??



- ❖ Shift from computation to information
  - at the “low end”: scramble to webspace (a mess!)
  - at the “high end”: scientific applications
- ❖ Datasets increasing in diversity and volume.
  - Digital libraries, interactive video, Human Genome project, EOS project
  - ... need for DBMS exploding
- ❖ DBMS encompasses most of CS
  - OS, languages, theory, “A”I, multimedia, logic

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

7

## Data Models




- ❖ A data model is a collection of concepts for describing data.
- ❖ A schema is a description of a particular collection of data, using the a given data model.
- ❖ The relational model of data is the most widely used model today.
  - Main concept: relation, basically a table with rows and columns.
  - Every relation has a schema, which describes the columns, or fields.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

8

## Levels of Abstraction



\* Schemas are defined using DDL; data is modified/queried using DML.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

9

## Example: University Database




- ❖ Conceptual schema:
  - Students(sid: string, name: string, login: string, dob: date, gpa:real)
  - Courses(cid: string, cname:string, credits:integer)
  - Enrolled(sid:string, cid:string, grade:string)
- ❖ Physical schema:
  - Relations stored as unordered files.
  - Index on first column of Students.
- ❖ External Schema (View):
  - Course\_info(cid:string,enrollment:integer)

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

10

## Relational Instance



Student table:

sid	name	login	dob	gpa
12345	Jones	jones@cs	1/1/1985	3.2
12346	Smith	smith@cs	2/3/1985	3.5
12347	Guldu	guldu@cs	3/4/1986	3.3

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

11

## Data Independence \*



- ❖ Applications insulated from how data is structured and stored.
- ❖ Logical data independence: Protection from changes in logical structure of data.
- ❖ Physical data independence: Protection from changes in physical structure of data.

\* One of the most important benefits of using a DBMS!

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

12

## Concurrency Control

- ❖ Concurrent execution of user programs is essential for good DBMS performance.
  - Because disk accesses are frequent, and relatively slow, it is important to keep the CPU humming by working on several user programs concurrently.
- ❖ Interleaving actions of different user programs can lead to inconsistency: e.g., check is cleared while account balance is being computed.
- ❖ DBMS ensures such problems don't arise: users can pretend they are using a single-user system.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

13

## Transaction: An Execution of a DB Program

- ❖ Key concept is transaction, which is an atomic sequence of database actions (reads/writes).
- ❖ Each transaction, executed completely, must leave the DB in a consistent state if DB is consistent when the transaction begins.
  - Users can specify some simple integrity constraints on the data, and the DBMS will enforce these constraints.
  - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
  - Thus, ensuring that a transaction (run alone) preserves consistency is ultimately the user's responsibility!

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

14

## Scheduling Concurrent Transactions

- ❖ DBMS ensures that execution of  $\{T_1, \dots, T_n\}$  is equivalent to some serial execution  $T'_1 \dots T'_n$ .
  - Before reading/writing an object, a transaction requests a lock on the object, and waits till the DBMS gives it the lock. All locks are released at the end of the transaction. (Strict 2PL locking protocol.)
  - **Idea:** If an action of  $T_i$  (say, writing X) affects  $T_j$  (which perhaps reads X), one of them, say  $T_i$ , will obtain the lock on X first and  $T_j$  is forced to wait until  $T_i$  completes; this effectively orders the transactions.
  - What if  $T_j$  already has a lock on Y and  $T_i$  later requests a lock on Y? (Deadlock)  $T_i$  or  $T_j$  is aborted and restarted!

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

15

## Ensuring Atomicity

- ❖ DBMS ensures atomicity (all-or-nothing property) even if system crashes in the middle of a Xact.
- ❖ **Idea:** Keep a log (history) of all actions carried out by the DBMS while executing a set of Xacts:
  - Before a change is made to the database, the corresponding log entry is forced to a safe location. (WAL protocol; OS support for this is often inadequate.)
  - After a crash, the effects of partially executed transactions are undo using the log. (Thanks to WAL, if log entry wasn't saved before the crash, corresponding change was not applied to database!)

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

16

## The Log

- ❖ The following actions are recorded in the log:
  - *Ti writes an object*: the old value and the new value.
    - Log record must go to disk before the changed page!
  - *Ti commits/aborts*: a log record indicating this action.
- ❖ Log records chained together by Xact id, so it's easy to undo a specific Xact (e.g., to resolve a deadlock).
- ❖ Log is often duplexed and archived on "stable" storage.
- ❖ All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

17

## Databases make these folks happy ...


- ❖ End users and DBMS vendors
  - ❖ DB application programmers
    - E.g. smart webmasters
  - ❖ Database administrator (DBA)
    - Designs logical / physical schemas
    - Handles security and authorization
    - Data availability, crash recovery
    - Database tuning as needs evolve
- Must understand how a DBMS works!*

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

18

## Structure of a DBMS

- ❖ A typical DBMS has a layered architecture.
- ❖ The figure does not show the concurrency control and recovery components.
- ❖ This is one of several possible architectures; each system has its own variations.



Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

19

## Summary

- ❖ DBMS used to maintain, query large datasets.
- ❖ Benefits include recovery from system crashes, concurrent access, quick application development, data integrity and security.
- ❖ Levels of abstraction give data independence.
- ❖ A DBMS typically has a layered architecture.
- ❖ DBAs hold responsible jobs and are **well-paid!**
- ❖ DBMS R&D is one of the broadest, most exciting areas in CS.



Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

20